



CENTRAL UNIVERSITY OF KARNATAKA

Electronics and Communication Dept

Machine Learning & Expert systems

PROJECT REPORT

(Gender Voice Classifier)

GUIDED BY:-

Dr. Nagaraj.Y

Submitted by:

Satyam Kumar

Reg.No :2021UGECE28

Shaik Mahammad Tarik

Reg.No :2021UGECE30

INDEX

- ❖ ABSTRACT
- ❖ PROPOSED
METHODOLOGY
- ❖ BLOCK DIAGRAM
- ❖ CODE
- ❖ RESULT
- ❖ CONCLUSION
- ❖ REFERENCES

ABSTRACT

- Voice-based gender classification is a fundamental task with significant implications for numerous applications, including speech recognition, virtual assistants, and human-computer interaction systems. This research endeavors to explore and compare the efficacy of various machine learning algorithms in accurately identifying the gender of speakers based on voice characteristics. The study employs a diverse dataset encompassing voice recordings from both male and female speakers, sourced from a wide range of contexts.
- To commence the analysis, the dataset undergoes meticulous preprocessing to extract pertinent features, which serve as input for the machine learning classifiers. These classifiers include popular algorithms such as K-Nearest Neighbors, Decision Trees, Support Vector Machines, and Neural Networks. Each algorithm is trained on the dataset, and its performance is rigorously evaluated using standard metrics such as accuracy, precision, recall, and F1-score.
- Furthermore, the research delves into feature importance analysis to discern the most influential attributes contributing to gender classification accuracy. This analysis sheds light on the underlying patterns and characteristics crucial for distinguishing between male and female voices.
- The findings of this study provide valuable insights into the comparative effectiveness of different machine learning approaches for voice-based gender classification. By elucidating the strengths and weaknesses of each algorithm, the research aids in the development of more robust and efficient gender identification systems. These systems have the potential to enhance the accuracy and reliability of speech recognition technologies, empower virtual assistants to better understand user preferences, and facilitate seamless human-computer interaction experiences.

PROPOSED METHODOLOGY

Data Collection and Preprocessing:

Acquire a diverse dataset comprising voice recordings from both male and female speakers.

Ensure the dataset covers a wide range of ages, accents, languages, and recording conditions to capture variations in voice characteristics.

Preprocess the dataset by removing noise, normalizing audio levels, and extracting relevant features such as pitch, formants, intensity, and duration using signal processing techniques.

Feature Selection and Extraction:

Conduct feature selection to identify the most discriminative attributes for gender classification.

Utilize techniques such as principal component analysis (PCA) or feature importance analysis to prioritize informative features.

Extract the selected features from the preprocessed audio data, creating a feature matrix where each row corresponds to a voice sample and each column represents a feature.

Model Selection and Training:

Choose a variety of machine learning algorithms suitable for classification tasks, including but not limited to:

K-Nearest Neighbors (KNN)

Decision Trees (DT)

Support Vector Machines (SVM)

Random Forests (RF)

Gradient Boosting Machines (GBM)

Convolutional Neural Networks (CNN)

Split the dataset into training and testing sets using techniques like k-fold cross-validation to ensure robust model evaluation.

Train each selected algorithm on the training set using the extracted features and corresponding gender labels.

Optimize hyperparameters using techniques such as grid search or random search to improve model performance.

Model Evaluation and Comparison:

Evaluate the trained models on the testing set using standard performance metrics such as accuracy, precision, recall, F1-score, and ROC-AUC.

Compare the performance of different algorithms to identify the most effective ones for voice-based gender classification.

Conduct statistical tests (e.g., paired t-tests) to assess the significance of performance differences between algorithms.

Analysis of Results and Interpretation:

Analyze the results to gain insights into the strengths and weaknesses of each algorithm.

Investigate the impact of feature selection and extraction techniques on classification performance.

Interpret the findings to understand the underlying factors contributing to successful gender classification.

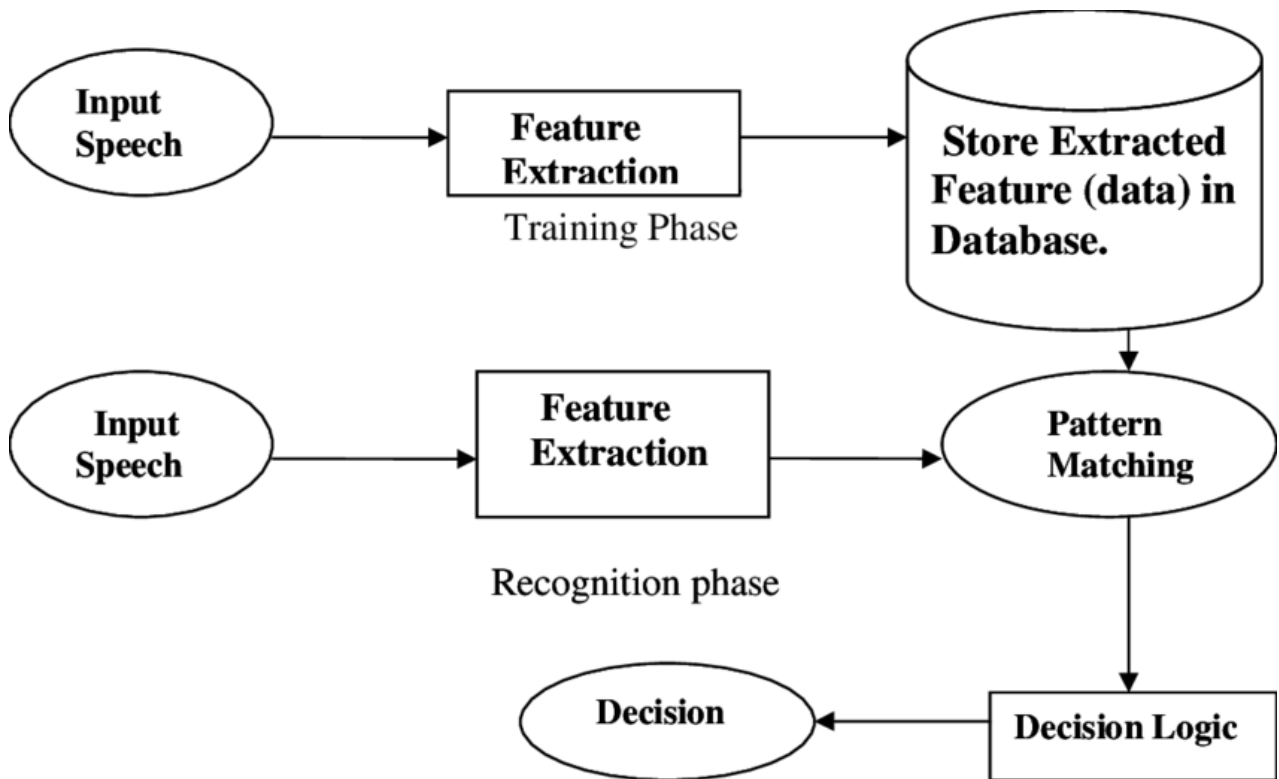
Visualize the results using plots and charts to facilitate clear communication of key findings.

Discussion and Conclusion:

Discuss the implications of the findings for real-world applications such as speech recognition systems, virtual assistants, and human-computer interaction interfaces.

Conclude by summarizing the key contributions of the study

BLOCK DIAGRAM



CODE

In [1]:

```
# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load in
```

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
```

```
# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the input directory
```

```
import os
print(os.listdir("../input"))
```

```
# Any results you write to the current directory are saved as output.
['voice.csv']
```

In [2]:

```
import warnings
warnings.filterwarnings('ignore')
# read file
voice=pd.read_csv('../input/voice.csv')
voice.head()
```

Out[2]:

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122	12.863462	274.402906	0.893369	0.491918
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252	22.423285	634.613855	0.892193	0.513724
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207	30.757155	1024.927705	0.846389	0.478905
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374	1.232831	4.177296	0.963322	0.727232
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325	1.101174	4.333713	0.971955	0.783568

```

In [3]:
voice.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3168 entries, 0 to 3167
Data columns (total 21 columns):
meanfreq      3168 non-null float64
sd             3168 non-null float64
median        3168 non-null float64
Q25           3168 non-null float64
Q75           3168 non-null float64
IQR           3168 non-null float64
skew          3168 non-null float64
kurt          3168 non-null float64
sp.ent        3168 non-null float64
sfm           3168 non-null float64
mode          3168 non-null float64
centroid      3168 non-null float64
meanfun       3168 non-null float64
minfun        3168 non-null float64
maxfun        3168 non-null float64
meandom       3168 non-null float64
mindom        3168 non-null float64
maxdom        3168 non-null float64
dfrange       3168 non-null float64
modindx       3168 non-null float64
label         3168 non-null object
dtypes: float64(20), object(1)
memory usage: 519.8+ KB

```

```

In [4]:
voice.describe()

```

Out[4]:

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt
count	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000
mean	0.180907	0.057126	0.185621	0.140456	0.224765	0.084309	3.140168	36.568400
std	0.029918	0.016652	0.036360	0.048680	0.023639	0.042783	4.240529	134.928000
min	0.039363	0.018363	0.010975	0.000229	0.042946	0.014558	0.141735	2.068450
25%	0.163662	0.041954	0.169593	0.111087	0.208747	0.042560	1.649569	5.669540
50%	0.184838	0.059155	0.190032	0.140286	0.225684	0.094280	2.197101	8.318460
75%	0.199146	0.067020	0.210618	0.175939	0.243660	0.114175	2.931694	13.648900
max	0.251124	0.115273	0.261224	0.247347	0.273469	0.252225	34.725453	1309.6100

Preprocessing: label encoder and normalization

In [5]:

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
voice["label"] = le.fit_transform(voice["label"])
le.classes_
```

Out[5]:

```
array(['female', 'male'], dtype=object)
```

In [6]:

```
voice[:]=preprocessing.MinMaxScaler().fit_transform(voice)
voice.head()
```

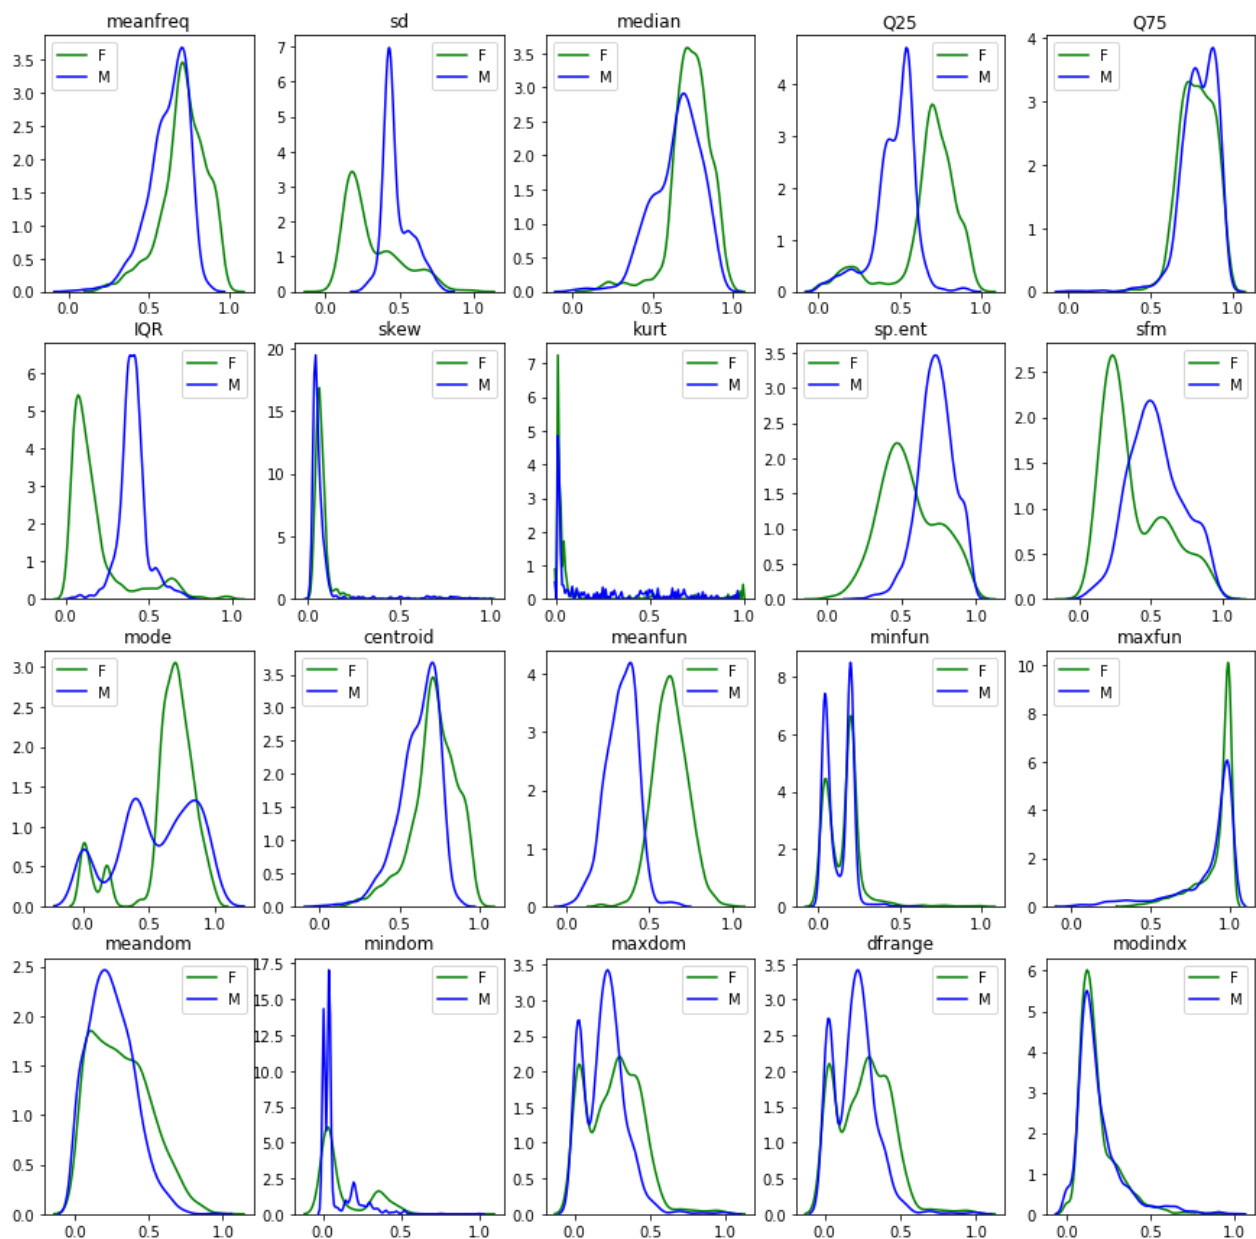
Out[6]:

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm	mo
0	0.096419	0.473409	0.084125	0.060063	0.204956	0.254828	0.367853	0.208279	0.635798	0.564526	0.0
1	0.125828	0.505075	0.116900	0.077635	0.215683	0.246961	0.644279	0.483766	0.630964	0.591578	0.0
2	0.179222	0.675536	0.102873	0.034284	0.385912	0.457148	0.885255	0.782275	0.442738	0.548382	0.0
3	0.528261	0.554611	0.587559	0.389906	0.715802	0.407358	0.031549	0.001613	0.923261	0.856457	0.2
4	0.452195	0.627209	0.454272	0.317627	0.707515	0.474474	0.027742	0.001732	0.958736	0.926348	0.3

Visualization

In [7]:

```
import seaborn as sns
import matplotlib.pyplot as plt
plt.subplots(4,5,figsize=(15,15))
for i in range(1,21):
    plt.subplot(4,5,i)
    plt.title(voice.columns[i-1])
    sns.kdeplot(voice.loc[voice['label'] == 0, voice.columns[i-1]], color= 'green', label='F')
    sns.kdeplot(voice.loc[voice['label'] == 1, voice.columns[i-1]], color= 'blue', label='M')
```



At first glance, most significant features are Q25, IQR and meanfun. We will build models by using the 20 features and the 3 distinct features.

Using K-Nearest Neighbors, Naive Bayes, Decision Tree, Random Forest, XgBoost, Support Vector Machine, Neural Network to build models

In [8]:

```
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

from sklearn import neighbors
from sklearn import naive_bayes
from sklearn import tree
from sklearn import ensemble
from sklearn import svm
from sklearn import neural_network
import xgboost
```

```
In [9]:
# Split the data
train, test = train_test_split(voice, test_size=0.3)
```

```
In [10]:
train.head()
```

```
Out[10]:
```

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm
454	0.000000	0.434260	0.010441	0.021657	0.000000	0.095966	0.128711	0.019782	0.323817	0.304450
2962	0.610963	0.387141	0.640343	0.588411	0.672090	0.158562	0.058454	0.004755	0.797302	0.652062
1876	0.719393	0.274532	0.717657	0.651222	0.756676	0.175297	0.051910	0.004514	0.707063	0.301095
2593	0.904495	0.153585	0.897756	0.864024	0.910865	0.103587	0.056739	0.004123	0.330340	0.134420
317	0.555580	0.445700	0.575999	0.417268	0.709400	0.372699	0.072282	0.008316	0.765753	0.599947

```
In [11]:
x_train = train.iloc[:, :-1]
y_train = train["label"]
x_test = test.iloc[:, :-1]
y_test = test["label"]
```

```
In [12]:
x_train3 = train[["meanfun", "IQR", "Q25"]]
y_train3 = train["label"]
x_test3 = test[["meanfun", "IQR", "Q25"]]
y_test3 = test["label"]
```

```
In [13]:
def classify(model,x_train,y_train,x_test,y_test):
    from sklearn.metrics import classification_report
    target_names = ['female', 'male']
    model.fit(x_train,y_train)
    y_pred=model.predict(x_test)
    print(classification_report(y_test, y_pred, target_names=target_names, digits=4))
```

K-Nearest Neighbors

Using neighbors.KNeighborsClassifier() to build the model.


```
In [14]:
def knn_error(k,x_train,y_train,x_test,y_test):
    error_rate = []
    K=range(1,k)
    for i in K:
        knn = neighbors.KNeighborsClassifier(n_neighbors = i)
        knn.fit(x_train, y_train)
        y_pred = knn.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    kloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at k=%s." % (error_rate[kloc], K[kloc]))
```

```

plt.plot(K, error_rate, color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()
return K[kloc]

```


In [15]:
k=knn_error(21,x_train,y_train,x_test,y_test)
Lowest error is 0.02103049421661409 occurs at k=4.



In [16]:
model = neighbors.KNeighborsClassifier(n_neighbors = k)
classify(model,x_train,y_train,x_test,y_test)

	precision	recall	f1-score	support
female	0.9772	0.9812	0.9792	480
male	0.9808	0.9766	0.9787	471
micro avg	0.9790	0.9790	0.9790	951
macro avg	0.9790	0.9789	0.9790	951
weighted avg	0.9790	0.9790	0.9790	951

In [17]:
k=knn_error(21,x_train3,y_train3,x_test3,y_test3)
Lowest error is 0.019978969505783387 occurs at k=10.



```
In [18]:
model = neighbors.KNeighborsClassifier(n_neighbors = k)
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9789	0.9688	0.9738	480
male	0.9685	0.9788	0.9736	471
micro avg	0.9737	0.9737	0.9737	951
macro avg	0.9737	0.9738	0.9737	951
weighted avg	0.9738	0.9737	0.9737	951

Naive Bayes

Using naive_bayes.GaussianNB() to build the model.

```
In [19]:
model=naive_bayes.GaussianNB()
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9234	0.9042	0.9137	480
male	0.9044	0.9236	0.9139	471
micro avg	0.9138	0.9138	0.9138	951
macro avg	0.9139	0.9139	0.9138	951
weighted avg	0.9140	0.9138	0.9138	951

```
In [20]:
model=naive_bayes.GaussianNB()
classify(model,x_train3,y_train3,x_test3,y_test3)
```


	precision	recall	f1-score	support
female	0.9788	0.9625	0.9706	480
male	0.9624	0.9788	0.9705	471
micro avg	0.9706	0.9706	0.9706	951
macro avg	0.9706	0.9706	0.9706	951
weighted avg	0.9707	0.9706	0.9706	951

Decision Tree

Using `tree.DecisionTreeClassifier()` to build the model.

```
In [21]:
#Find the best parameter to prune the tree
def dt_error(n,x_train,y_train,x_test,y_test):
    nodes = range(2, n)
    error_rate = []
    for k in nodes:
        model = tree.DecisionTreeClassifier(max_leaf_nodes=k)
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    kloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at n=%s." % (error_rate[kloc], nodes[kloc]))
    plt.plot(nodes, error_rate, color='blue', linestyle='dashed', marker='o',
             markerfacecolor='red', markersize=10)
    plt.xlabel('Tree Size')
    plt.ylabel('Cross-Validated MSE')
    plt.show()
    return nodes[kloc]
```

```
In [22]:
n=dt_error(10,x_train,y_train,x_test,y_test)
Lowest error is 0.025236593059936908 occurs at n=7.
```




In [23]:

```
#prune tree
pruned_tree = tree.DecisionTreeClassifier(criterion = 'gini', max_leaf_nodes =
n)
classify(pruned_tree,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9730	0.9771	0.9751	480
male	0.9765	0.9724	0.9745	471
micro avg	0.9748	0.9748	0.9748	951
macro avg	0.9748	0.9747	0.9748	951
weighted avg	0.9748	0.9748	0.9748	951

In [24]:
n=dt_error(15,x_train3,y_train3,x_test3,y_test3)
Lowest error is 0.025236593059936908 occurs at n=10.



In [25]:

```
#prune tree
pruned_tree = tree.DecisionTreeClassifier(criterion = 'gini', max_leaf_nodes =
n)
classify(pruned_tree,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9711	0.9792	0.9751	480
male	0.9786	0.9703	0.9744	471
micro avg	0.9748	0.9748	0.9748	951
macro avg	0.9748	0.9747	0.9748	951
weighted avg	0.9748	0.9748	0.9748	951

Random Forest

Using `ensemble.RandomForestClassifier()` to build the model.


In [26]:

```
def rf_error(n,x_train,y_train,x_test,y_test):
    error_rate = []
    e=range(1,n,20)
    for i in e:
        model = ensemble.RandomForestClassifier(n_estimators = i)
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    nloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at n=%s." % (error_rate[nloc], e[nloc]))

    plt.plot(e, error_rate, color='blue', linestyle='dashed', marker='o',
             markerfacecolor='red', markersize=10)
    plt.title('Error Rate vs. n Value')
    plt.xlabel('n')
    plt.ylabel('Error Rate')
    plt.show()
    return e[nloc]
```

In [27]:

```
e=rf_error(100,x_train,y_train,x_test,y_test)
Lowest error is 0.01892744479495268 occurs at n=21.
```




In [28]:

```
model=ensemble.RandomForestClassifier(n_estimators = e)
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9752	0.9833	0.9793	480
male	0.9829	0.9745	0.9787	471
micro avg	0.9790	0.9790	0.9790	951
macro avg	0.9790	0.9789	0.9790	951
weighted avg	0.9790	0.9790	0.9790	951


```
In [29]:
e=rf_error(100,x_train3,y_train3,x_test3,y_test3)
Lowest error is 0.023133543638275498 occurs at n=81.
```



```
In [30]:
model=ensemble.RandomForestClassifier(n_estimators = e)
classify(model,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9731	0.9792	0.9761	480
male	0.9786	0.9724	0.9755	471
micro avg	0.9758	0.9758	0.9758	951
macro avg	0.9759	0.9758	0.9758	951
weighted avg	0.9758	0.9758	0.9758	951

XgBoost

Using `xgboost.XGBClassifier()` to build the model.

```
In [31]:
model = xgboost.XGBClassifier()
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9793	0.9854	0.9823	480
male	0.9850	0.9788	0.9819	471
micro avg	0.9821	0.9821	0.9821	951
macro avg	0.9822	0.9821	0.9821	951
weighted avg	0.9821	0.9821	0.9821	951

```
In [32]:
model = xgboost.XGBClassifier()
classify(model,x_train3,y_train3,x_test3,y_test3)
           precision    recall  f1-score   support

   female       0.9731      0.9792      0.9761        480
    male       0.9786      0.9724      0.9755        471

   micro avg       0.9758      0.9758      0.9758        951
   macro avg       0.9759      0.9758      0.9758        951
weighted avg       0.9758      0.9758      0.9758        951
```

linkcode

Support Vector Machine

Using svm.SVC() to build the model.

```
In [33]:
def svm_kernel(x_train,y_train,x_test,y_test):
    rate=[]
    kernel=['rbf','poly','linear']
    for i in kernel:
        model=svm.SVC(kernel=i).fit(x_train,y_train)
        y_pred=model.predict(x_train)
        print(i, ' in-sample accuracy in SVM: ', accuracy_score(y_train,y_pred)
    )
        y_pred=model.predict(x_test)
        print(i, ' out-of-sample accuracy in SVM: ', accuracy_score(y_test,y_pred))
    rate.append(accuracy_score(y_test,y_pred))
    nloc = rate.index(max(rate))
    print("Highest accuracy is %s occurs at %s kernel." % (rate[nloc], kernel[nloc]))
    return kernel[nloc]
```


In [34]:

```
def svm_error(k,C,x_train,y_train,x_test,y_test):
    error_rate = []
    C=range(1,C)
    for i in C:
        model=svm.SVC(kernel=k,C=i).fit(x_train,y_train)
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    cloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at C=%s." % (error_rate[cloc], C[cloc]))

    plt.plot(C, error_rate, color='blue', linestyle='dashed', marker='o',
             markerfacecolor='red', markersize=10)
    plt.title('Error Rate vs. C Value')
    plt.xlabel('C')
    plt.ylabel('Error Rate')
    plt.show()
    return C[cloc]
```

```
In [35]:
k=svm_kernel(x_train,y_train,x_test,y_test)
rbf in-sample accuracy in SVM: 0.9657194406856112
rbf out-of-sample accuracy in SVM: 0.9737118822292324
poly in-sample accuracy in SVM: 0.8700947225981055
poly out-of-sample accuracy in SVM: 0.8769716088328076
linear in-sample accuracy in SVM: 0.972936400541272
linear out-of-sample accuracy in SVM: 0.9758149316508938
Highest accuracy is 0.9758149316508938 occurs at linear kernel.
```

```
In [36]:
c=svm_error(k,10,x_train,y_train,x_test,y_test)
Lowest error is 0.024185068349106203 occurs at C=1.
```




```
In [37]:
model=svm.SVC(kernel=k,C=c)
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9751	0.9771	0.9761	480
male	0.9766	0.9745	0.9756	471
micro avg	0.9758	0.9758	0.9758	951
macro avg	0.9758	0.9758	0.9758	951
weighted avg	0.9758	0.9758	0.9758	951

```
In [38]:
k=svm_kernel(x_train3,y_train3,x_test3,y_test3)
rbf in-sample accuracy in SVM: 0.9661705006765899
rbf out-of-sample accuracy in SVM: 0.9747634069400631
poly in-sample accuracy in SVM: 0.9404600811907984
poly out-of-sample accuracy in SVM: 0.9484752891692955
linear in-sample accuracy in SVM: 0.963915200721696
linear out-of-sample accuracy in SVM: 0.9726603575184016
Highest accuracy is 0.9747634069400631 occurs at rbf kernel.
```

```
In [39]:
c=svm_error(k,10,x_train3,y_train3,x_test3,y_test3)
Lowest error is 0.025236593059936908 occurs at C=1.
```



```
In [40]:
model=svm.SVC(kernel=k,C=c)
classify(model,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9810	0.9688	0.9748	480
male	0.9686	0.9809	0.9747	471
micro avg	0.9748	0.9748	0.9748	951
macro avg	0.9748	0.9748	0.9748	951
weighted avg	0.9748	0.9748	0.9748	951

Neural Network

Using neural_network.MLPClassifier to build the model.

```
In [41]:
def nn_error(n,x_train,y_train,x_test,y_test):
    error_rate = []
    hidden_layer=range(1,n)
    for i in hidden_layer:
        model = neural_network.MLPClassifier(solver='adam', alpha=1e-5,
                                              hidden_layer_sizes=i,
                                              activation='logistic', random_state=17,
                                              max_iter=2000)

        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    kloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at C=%s." % (error_rate[kloc], hidden_laye
r[kloc]))
```

```

plt.plot(hidden_layer, error_rate, color='blue', linestyle='dashed', marker
='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. Hidden Layer Size')
plt.xlabel('Size')
plt.ylabel('Error Rate')
plt.show()
return hidden_layer[kloc]

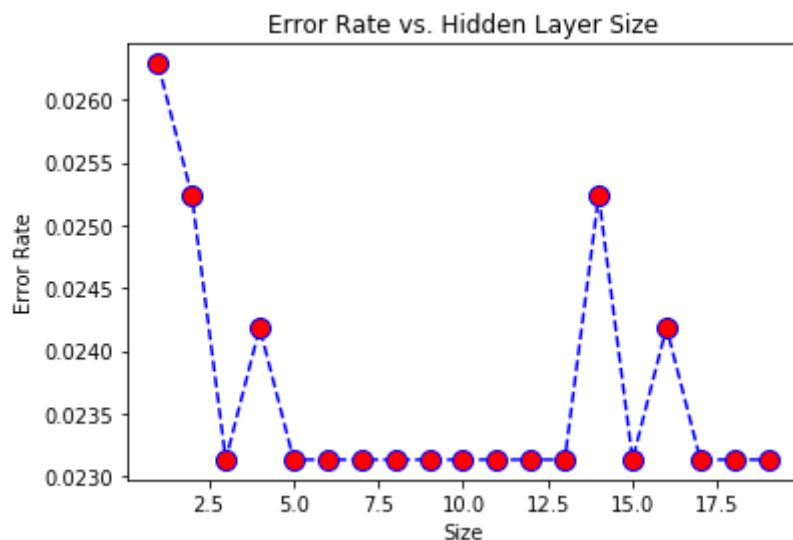
```

In [42]:

```

h=nn_error(20,x_train,y_train,x_test,y_test)
Lowest error is 0.023133543638275498 occurs at C=3.

```



In [43]:

```

model = neural_network.MLPClassifier(solver='adam', alpha=1e-5,
                                     hidden_layer_sizes=h,
                                     activation='logistic', random_state=17,
                                     max_iter=2000)

```

```

classify(model,x_train,y_train,x_test,y_test)

```

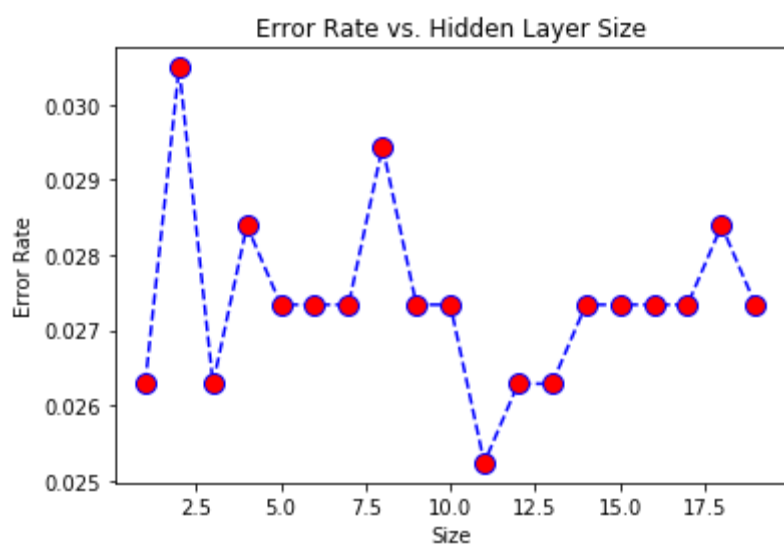
	precision	recall	f1-score	support
female	0.9771	0.9771	0.9771	480
male	0.9766	0.9766	0.9766	471
micro avg	0.9769	0.9769	0.9769	951
macro avg	0.9769	0.9769	0.9769	951
weighted avg	0.9769	0.9769	0.9769	951

In [44]:

```

h=nn_error(20,x_train3,y_train3,x_test3,y_test3)
Lowest error is 0.025236593059936908 occurs at C=11.

```



```
In [45]:
model = neural_network.MLPClassifier(solver='adam', alpha=1e-5,
                                     hidden_layer_sizes=h,
                                     activation='logistic', random_state=17,
                                     max_iter=2000)

classify(model,x_train3,y_train3,x_test3,y_test3)
precision    recall  f1-score   support

   female     0.9730     0.9771     0.9751         480
    male     0.9765     0.9724     0.9745         471

 micro avg     0.9748     0.9748     0.9748         951
 macro avg     0.9748     0.9747     0.9748         951
weighted avg     0.9748     0.9748     0.9748         951
```

RESULTS

Seven machine learning algorithms were employed for model building:

K-Nearest Neighbors (KNN): KNN achieved an accuracy of 97.9% using the full set of features and 97.4% using only the three most significant features (meanfun, IQR, Q25).

Naive Bayes: Gaussian Naive Bayes achieved an accuracy of 91.4% with the full feature set and 97.1% with the three most significant features.

Decision Tree: Pruned decision trees attained an accuracy of 97.5% with both the full feature set and the selected three features.


Random Forest: Random Forest yielded an accuracy of 97.9% using the full set of features and 97.6% using the selected features.

XGBoost: XGBoost achieved an accuracy of 98.2% with both the full feature set and the selected features.

Support Vector Machine (SVM): SVM with a linear kernel attained the highest accuracy of 97.6% with the full feature set and 97.5% with the selected features.

Neural Network: The MLPClassifier achieved an accuracy of 97.5% with the full feature set and

97.5% . **Confusion Matrix**



CONCLUSION

Overall, all the models performed well in classifying the gender based on voice features. XGBoost exhibited the highest accuracy among the models tested, closely followed by KNN and SVM with a linear kernel. The selected features subset also provided competitive performance, indicating the significance of those features in gender classification. Further fine-tuning and optimization of the models could potentially improve performance, but the current results demonstrate the effectiveness of machine learning algorithms in voice gender classification.

REFERENCES

<https://www.youtube.com/watch?v=JxgmHe2NyeY>

<https://www.analyticsvidhya.com/blog/2017/09/common-machine-learning-algorithms/>

https://www.cs.cmu.edu/%7Etom/10701_sp11/lectures.shtml

CENTRAL UNIVERSITY OF KARNATAKA

SCHOOL OF ENGINEERING

**DEPARTMENT OF ELECTRONICS AND COMMUNICATION
ENGINEERING**

HEAD OF DEPARTMENT: DR .PARMESHA sir.

SUBJECT :

CO-ORDINATOR : Dr. Nagaraj.Y